

OmMeGo.AI Mission: Super-Intelligence for your Organization. At your fingertips.

Signup: contact business@ommegeo.ai to create your own org id!

Apache 2.0 Github release: click [here](#).

Note: Apache 2.0 license release includes paper, slides meant as a tutorial. Recipe and templates are available after you sign up, but some subset will be released as Apache 2.0 soon



Getting setup with [OmMeGo.AI](https://ommegogo.ai) & how to use it

Initial setup of [OmMeGo.AI](https://ommegogo.ai) super-intelligence for your organization

Setup

1. Send email request to business@ommegogo.ai with your org details.
2. We create a org-id for you (e.g. xyz)
3. Business manager signs up at [OmMeGo.AI](https://ommegogo.ai) with “xyz” org id.
4. Manager approves team members: additional managers, agent creators, agent owners and agent users from your org.
5. Team member domain experts build automation using new sub-agents, hosted privately in your org @ [OmMeGo.AI](https://ommegogo.ai). No one outside your org can see them.

Usage

6. Instant org-level access to all team-members the combined super-intelligence of all sub-agents using our breakthrough agent long-term memory tech (router).
7. Free powerful templates for recipes and higher-level sub-agents from [OmMeGo.AI](https://ommegogo.ai) enable RSI (recursive self-improvement) of your sub-agents.
8. **[OmMeGo.AI](https://ommegogo.ai) Certified Agent Developers**: train your team-members to build your own RSI recipes and sub-agents. Can also build them and manage them for you.



Target Audience: who want process automation, but...

... don't want to spend time and money for custom software that will break often, as processes change!

1. Domain Experts (DE)

- Small DE businesses: independent clinics, coaches, pharmacies, accountants, real-estate agents, grocery stores, and more.
- Medium DE businesses: hospitals, banks, manufacturers.
- Large DE businesses: large software and hardware product companies.

2. Agentic Developer Consultants (AD)

- Get certified with [OmMeGo.AI](#).
- Monetize reusable recipe and agents for your customers.
- Help DE achieve 10x to 100x in business process:
 - speed, accuracy, team collaboration
 - reducing mistakes, SLA, costs and wastes
 - many new savings and monetization not possible now



Platform Fees

What you pay when you signup

- **Agent IDE (not billed to OmMeGo.AI)**
 - Cursor is required: 20\$/user*. Pay [Cursor](#) directly**.
 - Other IDEs can be supported by our ADs if you use another one already.
- **[OmMeGo.AI](#) subscription fees**
 - 40\$ / month / user
 - Special deal: 1 month free trial, then 20\$ for the next 5 months.
 - Unlimited platform support included
- **[OmMeGo.AI](#) Certified AD fees**
 - Unlimited platform support included
 - Cost of building agents: we give free 5 hours of training & training material for your DE
 - Option 1 (for larger companies): hire a certified AD to work with your DE.
 - We don't make money on AD fees: you pay them directly.
 - We provide ADs training and enforce quality control in what they b
 - Option 2: we provide free templates and tutorials to your DE to do it inhouse.


** If you use cursor a lot, you may have to pay higher fees, but Cursor is a lot cheaper than the human labor time!*

***[OmMeGo.com](#) and [OmMeGo.AI](#) developers use Cursor internally but Cursor is not directly associated with our company.*



OmMeGo.com: Hybrid Agents for Customer Coaching

Human-AI Hybrids for doctors, coaches etc. **High interactions human <-> agent <-> expert. Limited between agent <-> agent**




The synergy between the Clone Coaches and our Human coaches is seamless. For straightforward inquiries, the clone provides quick, effective advice, only escalating complex issues to human coaches. This efficient triage not only improves member satisfaction but also boosts our coaching productivity.



OmMeGo.AI: Super-agents for organizational Super-intelligence

Lots of sub-agents with High interactions between agent <-> agent and agent <-> system. Limited between human <-> agent



OmMeGo.AI: Soup Architecture

Goal: 10-100x org-level complex processes

- Highly interacting agents bounded by the Router
- Unified by the Long-Term Memory Router
- Automatically generated for each Org
- Router Self-learns and adapts over time
- Learns and uses automatic linkages between sub-agents

OmMeGo.com: Sunny-side Architecture

Goal: 10-100x human experts <-> customer & expert <-> expert

- Customers want to talk to one expert at a time
- Interactions between experts handled by their agents
- Agent clone of the human expert wraps around the human and fields all the chats from human customers



OmMeGo.AI SuperAgent Platform: Core Concepts

AD

Key terms to understand what they do. Topical routers is a key innovation from [OmMeGo.AI](#) that makes everything work

Concept	Description
Subagent	Agentic brain with the full knowledge of a sub-system: current architecture, future improvements coming soon. Also contains end-to-end recipes that can do work on this sub-systems such as run tools, APIs. You can think of Anthropic skills as a subset of an OmMeGo.AI sub-agent, which has a lot more structure and powers
Recipe	An agentic computer program. It can do a set of well-defined things predictably, and unlike regular computer programs, it can also perform a large number of extensions of the basic tasks it was never programmed for.
Routers	A collection of look-up table built using your org sub-agents brain + some free templates, agents and recipes we give to you. Routers are learned automatically using a custom algorithm developed by OmMeGo.AI (discussed in the paper), but can also be updated incrementally by org managers.



Concept	Description
Superagent	Routers when installed properly into your Agentic IDE glue your sub-agents together and give the IDE high-quality long-term memory of ALL of your business knowledge encoded in your sub-agents.
Agent RSI	Your custom agents when hosted on OmMeGo.AI come with ability to get better over time automatically in a few different ways: • Sub-agent brain get updated by users after a new task finishes not done before by user prompt • Recipes can be updated when they accomplish new task variation by simple user update, or if enabled can self-update automatically • Router keeps auto-learning about your new agents and their powers as they get updated • Your higher level templates help your system know how to learn: updated by your managers based on feedbacks.



Software Development Ladder: before modern Agents

AD

A "ladder" of bi-directional vision-based communication was used: developer <-> system


Supervision* was the key:

Human developer needs to know what's going on in order to build the system. Development gives continuous live feedback making it fast and efficient:

- IDEs** provided Human with Vision
- Tests provide System with Vision

**ML term when labeled data is shown to an intelligent system to help it learn*

***Integrated Development Environment*






Natural progression to modern Agents: 3-D ladder system

AD

[OmMeGo.AI](#) paper and platform introduces a formal 3-d ladder Agent vision system, natural descendant of the 2-d system

- SuperVision*: Supervision key for BOTH Agent & Human for learning/reliable development
 - UX read by Super Agent
 - Agentic IDE* for design
 - Test cases reviewed by Super Agent
- SuperAgent: router index glues all sub-agents together
- Bi-directional 3-D ladder communication:
 - Human<->Agent (HA)
 - Agent<->System (AS)
 - Human<->System (HS)



*Cursor, Antigravity, OpenCode



OmMeGo.AI The Agentic Work Ladder: L1 to L5

AD

OmMeGo.AI introduces multi-level agents: from faster design and coding to self-improving agent systems

Old ladder: human developer ↔ system feedback

New ladder: user ↔ agent ↔ agent-authoring ↔ router-authoring ↔ meta-learning

L1 — Creative direct use

L2 — Reliable known-task execution

L3 — Agent/recipe authoring

L4 — Agent template structure authoring

L5 — Router, other meta-algorithm authoring

- The same person may operate at different levels at different times.
- Most end-user work lives in **L1 / L2**.
- Organizational learning and **RSI** happen through **L3 / L4 / L5**.

Coming soon: L0 — Minion layer! Execution helpers for narrowly-scoped delegated tasks.

Not a pricing tier. Not a job title. A depth-of-authoring model for how work moves from using an agent to improving the agent ecosystem itself.



Work Ladder deep-dive: L1 — Creative sub-agent usage

AD

Work ladder deep-dive; how layers interact with each other

- Direct use for problems that are **not fully predefined**.
- This is the creative mode: the agent must combine its existing brain with rich user context to solve something slightly new.
- It can be powerful but less stable, and may require multiple prompts.



Work Ladder deep-dive: L2 — Reliable sub-agent usage



Work ladder deep-dive; how layers interact with each other

- Direct use for complex but already well-understood tasks that the sub-agent brain already knows how to solve.
- This is the dominant day-to-day use case: short prompts, long useful answers, with no change to the core brain.



Work Ladder deep-dive: L3 — Sub-agent authoring

AD

Work ladder deep-dive; how layers interact with each other

- Indirect use where successful L1 discoveries — from human input or agent vision/self-discovery — are converted into new reusable knowledge by writing or updating sub-agents and recipes.
- This is where the system starts to improve itself in a durable way.



Work Ladder deep-dive: L4 — structure access & maint.

AD

Work ladder deep-dive; how layers interact with each other

- Automatically maintained and/or your org manager/architect owns it.
- Generate topic indexes, context bundles, and routing structure using the router builder.
- Sub-agent templates on how recipe “brains” are organized.
- Sub-agent templates on how your sub-agent brain are organized.
- Instructions on how the agent cross-references are maintained so that org-wide knowledge graph can be maintained.
- Agentic brain merge algorithms



Work Ladder deep-dive: L5 — Meta-algorithm authoring

AD

Work ladder deep-dive; how layers interact with each other


- Typically maintained by [OmMeGo.AI](#) (or your org for self deployment)
- Updating level 4 agent and recipe instructions
- How the topic discovery algorithm works on your sub-agents
- How the router indexing algorithm works
- How agentic brain updates and merges work
- How the memory refresh works using the router
- Layered access controls provided by [OmMeGo.AI](#) to creators, owners, developers and users of sub-agents allowing the structures to not decay and get corrupted over time

L1/L2 solve tasks. L3/L4/L5 turn successful task-solving into reusable organizational intelligence.



Feedback self-learning in RSI recipes

Closed-loop improvement of organization-owned artifacts – note model fine-tuning



Rapid in-session improvement possible by recursively updating the brain after each step. Knows if the output improved by looking at it – producing the Agent (Super)Vision described earlier. Defects decrease as recipe matures learning on its own!



OmMeGo.AI Why is this Agentic Work ladder innovative?

AD

Why layered agents are different from ordinary AI tooling

Ordinary AI Tools	OmMeGo layered model	Result
• Mostly answer user prompts • Little Durable Improvement • Knowledge Stays in the session • Human has to discover solutions repeatedly	• L1/L2 solve work now • L3 converts successful new work into reusable agent knowledge • L4 organizes that knowledge into router structures • L5 improves how the system itself learns and refreshes memory	• Agent systems can grow from usage • Good outputs can become reusable structures • Organizations accumulate know-how instead of losing it • This is a path towards RSI / ASI-like self improvement loops

From answering tasks → to capturing knowledge → to improving routing → to improving learning itself